


Artificial Intelligence

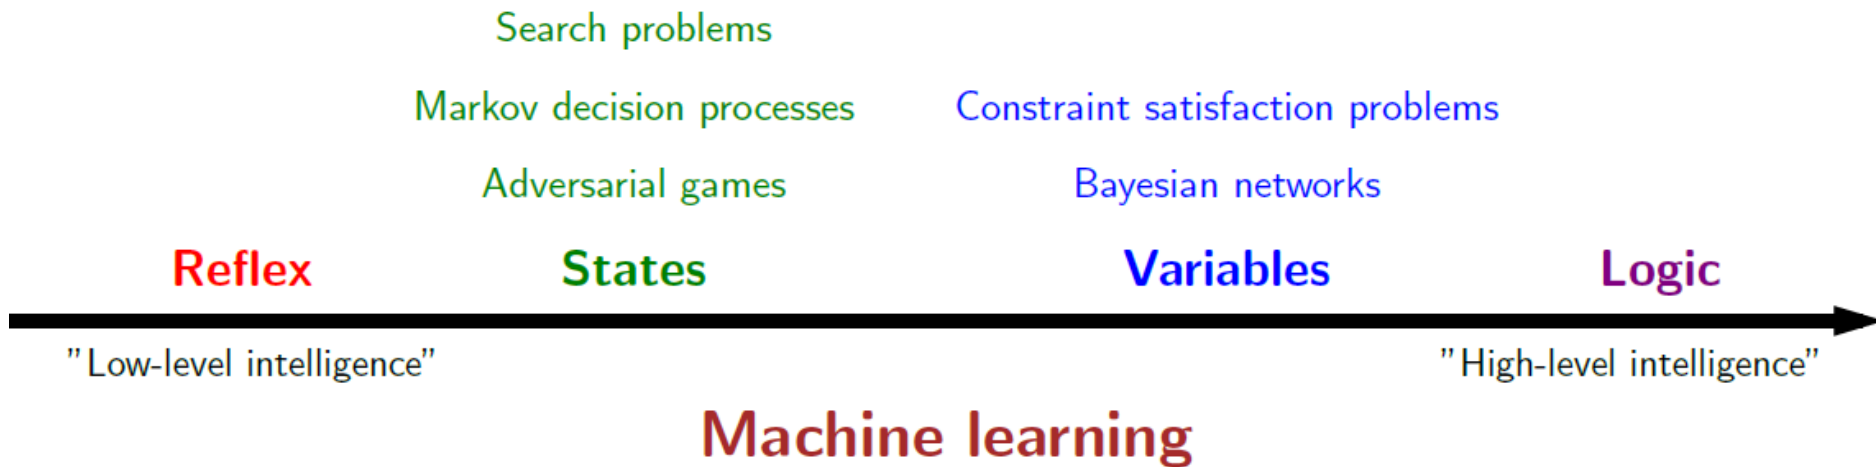
Adversarial Search




Logistics

- Midterm
 - Nov. 2nd, in class test
 - All questions will be related/similar to quiz/homework
- Final project
 - Will be released soon
- Final exam examples
 - Canvas

Course Topics




Graph Search



State Space Graphs



- State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



World states? $120 \times (2^{30}) \times (12^2) \times 4$

Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.



Graph Search

- Idea: never **expand** a state twice
- How to implement:
 - Tree search + set of expanded states (“closed set”)
 - Expand the search tree node-by-node, but...
 - Before **expanding a node**, check to make sure its state has **never been expanded before**
 - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list





Tree/ Graph Search Pseudo-Code

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
    end
  end
```

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
    end
  end
```



Consistency of Heuristics

- Consequences of consistency:
 - The f value along a path never decreases
 - $$h(A) \leq \text{cost}(A \text{ to } C) + h(C)$$
 - A* graph search is optimal




A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems




Optimality

- Tree search:
 - A* is optimal if heuristic is **admissible**
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is **consistent**
 - UCS optimal ($h = 0$ is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from **relaxed problems**




Adversarial Games




Types of Games

- Many different kinds of games!
- Axes:
 - Deterministic or stochastic?
 - One, two, or more players?
 - Zero sum?
 - Perfect information (can you see the state)?




Game Playing State-of-the-Art

- **Checkers:** 1950: First computer player. 1994: First computer champion: Chinook ended 40-year-reign of human champion Marion Tinsley using complete 8-piece endgame. 2007: Checkers solved!
- **Chess:** 1997: Deep Blue defeats human champion Gary Kasparov in a six-game match.
- **Go:** 2016/2017: AlphaGo defeats Korean and Chinese champions with reinforcement learning + deep learning.
- **Pacman**




Behavior from Computation




[Demo: mystery pacman (L6D1)]

Video of Demo Mystery Pacman




Rational Agent: Maximize Your Expected Utility




Act rationally: Want algorithms for calculating a **strategy (policy)** which recommends a move from each state


Deterministic Games

- Many possible formalizations, one is:
 - States: S (start at s_0)
 - Players: $P=\{1\dots N\}$ (usually take turns)
 - Actions: A (may depend on player / state)
 - Transition Function: $S \times A \rightarrow S$
 - Terminal Test: $S \rightarrow \{\text{true}, \text{false}\}$
 - Terminal Utilities: $S \times P \rightarrow R$
- Solution for a player is a **policy**: $S \rightarrow A$






Start State




Goal State

Single-Agent Trees



Value of a State


Value of a state:
The best
achievable
outcome (utility)
from that state



Non-Terminal States:
$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

Terminal States:
 $V(s) = \text{known}$

Zero-Sum Games




- Zero-Sum Games

- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition


- General Games

- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible
- More later on non-zero-sum games

Adversarial Search



Adversarial Game Trees




Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:


$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:


$$V(s) = \text{known}$$

Tic-Tac-Toe Game Tree




Adversarial Search (Minimax)

- Deterministic, zero-sum games:
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- Minimax search:
 - A state-space search tree
 - Players alternate turns
 - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



Minimax Implementation

Minimax values:
computed recursively



Terminal values:
part of the game


def max-value(s):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

return v



initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

return v

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Implementation (Dispatch)

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)


def max-value(state):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v




initialize $v = +\infty$

for each successor of state:




$v = \min(v, \text{value}(\text{successor}))$

return v

Minimax Example




Minimax Properties




Optimal against a perfect player. Otherwise?

[Demo: min vs exp (L6D2, L6D3)]


Video of Demo Min vs. Exp (Min)




Video of Demo Min vs. Exp (Exp)



Video of Demo Smart Ghosts (Coordination)




Video of Demo Smart Ghosts (Coordination) – Zoomed In



Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?




Resource Limits




Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
 - Instead, search only to a **limited depth** in the tree
 - Replace terminal utilities with an **evaluation function for non-terminal positions**
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm




Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation




[Demo: depth limited (L6D4, L6D5)]


Video of Demo Limited Depth (2)



Video of Demo Limited Depth (10)




Evaluation Functions



Evaluation Functions


- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:
 - e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

Evaluation for Pacman




[Demo: thrashing d=2, thrashing d=2 (fixed evaluation function), smart ghosts coordinate (L6D6,7,8,10)]


Game Tree Pruning



Minimax Example




Minimax Pruning



Alpha-Beta Pruning

- General configuration (MIN version)
 - We're computing the MIN-VALUE at some **node n**
 - We're looping over **n 's children**
 - n 's estimate of the childrens' min is dropping
 - Who cares about n 's value? **MAX**
 - Let **a** be the best value that **MAX** can get at any choice point along the current path from the root
 - If **n** becomes worse than **a** , MAX will avoid it, so we can stop considering **n** 's other children (it's already bad enough that it won't be played)
- MAX version is symmetric



Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root


MAX

MIN

⋮

MAX

MIN



initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \geq \beta$ return v

$\alpha = \max(\alpha, v)$

return v

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


if $v \leq \alpha$ return v

$\beta = \min(\beta, v)$

return v

Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of, e.g. chess, is still hopeless...
- Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - So the most naïve version won't let you do action selection
- This is a simple example of **metareasoning** (computing about what to compute)



Course Topics

